

Parallel Processing Fundamentals



Lesson Objectives

After completing this lesson, students should be able to:

- ◆ Describe how MapReduce works
 - Explain the reliance on the Key Value Pair (KVP) paradigm
 - Illustrate the MapReduce framework with simple examples

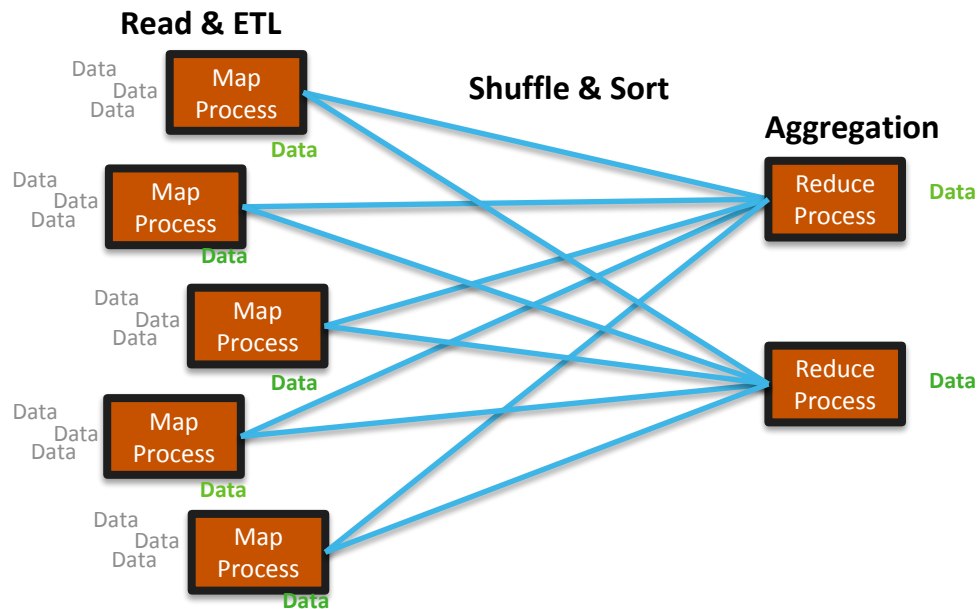


The MapReduce Framework



What is MapReduce?

Breaking a large problem into sub-solutions



Simple Algorithm

1. Review stack of quarters
2. Count each year that ends in an even number



Processing at Scale




Distributed Algorithm – MapReduce

Map
(total number of quarters)



Reduce
(sum each person's total)



The Mapper

- ◆ **The Mapper reads data in the form of key/value pairs (KVPs)**
- ◆ **It outputs zero or more KVPs**
- ◆ **The Mapper may use or completely ignore the input key**
 - For example, a standard pattern is to read a line of a file at a time
 - The key is the byte offset into the file at which the line starts
 - The value is the contents of the line itself
 - Typically the key is considered irrelevant with this pattern
- ◆ **If the Mapper writes anything out, it must in the form of KVPs**
 - This “intermediate data” is NOT stored in HDFS (local storage only without replication)

MapReduce Example – Map Phase

Input to Mapper

```
(8675, 'I will not eat green eggs and ham')  
(8709, 'I will not eat them Sam I am')  
...
```

- Ignoring the key
 - It is just an offset

Output from Mapper

```
('I', 1), ('will', 1), ('not', 1), ('eat', 1),  
('green', 1),  
('eggs', 1), ('and', 1),  
('ham', 1), ('I', 1), ('will', 1),  
('not', 1), ('eat', 1),  
('them', 1), ('Sam', 1),  
('I', 1), ('am', 1)
```

- In this example
 - The size of the output > size of the input
 - No attempt is made to optimize within a record in this example
 - This is a great use case for a “Combiner”



The Shuffle

- After the Map phase is over, all the outputs from the mappers are sent to reducers
- KVPs with the same key will be sent to the same reducer
 - By default (k,v) will be sent to the reducer number $\text{hash}(k) \% \text{numReducers}$
- This can potentially generate a lot of network traffic on your cluster**
 - In our word count example the size of the output data is of the same order of magnitude as our input data
- Some very common operations like join, or group by require a lot of shuffle by design
- Optimizing these operations is an important part of mastering distributed processing programming
- CPU and RAM can scale by adding worker nodes, network can't



MapReduce Example – Reduce Phase

Input to Reducer

```
('I', [1, 1, 1])
('Sam', [1])
('am', [1])
('and', [1])
('eat', [1, 1])
('eggs', [1])
('green', [1])
('ham', [1])
('not', [1, 1])
('them', [1])
('will', [1, 1])
```

Output from Reducer

```
('I', 3)
('Sam', 1)
('am', 1)
('and', 1)
('eat', 2)
('eggs', 1)
('green', 1)
('ham', 1)
('not', 2)
('them', 1)
('will', 2)
```

- Notice keys are sorted and associated values for same key are in a single list
 - Shuffle & Sort did this for us

- All done!

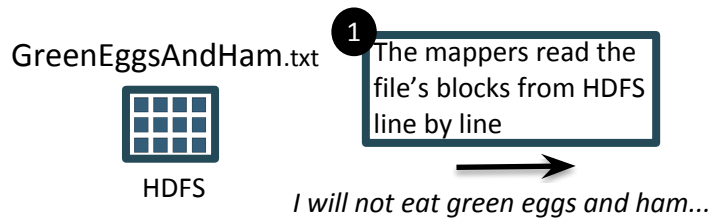


The Reducer

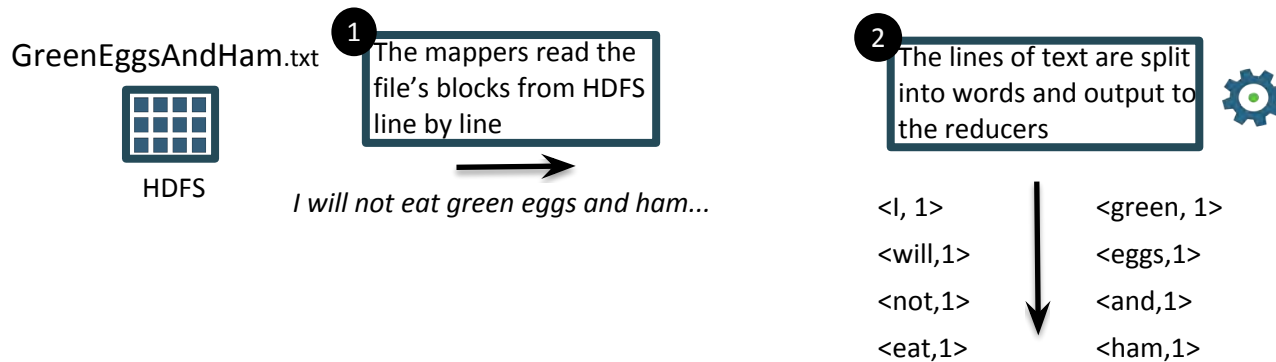
- ◆ **After the Shuffle phase is over, all the intermediate values for a given intermediate key are sorted and combined together into a list**
- ◆ **This list is given to a Reducer**
 - There may be a single Reducer, or multiple Reducers
 - All values associated with a particular intermediate key are guaranteed to go to the same Reducer
 - The intermediate keys, and their value lists, are passed in sorted order
- ◆ **The Reducer outputs zero or more KVPs**
 - These are written to HDFS
 - In practice, the Reducer often emits a single KVP for each input key



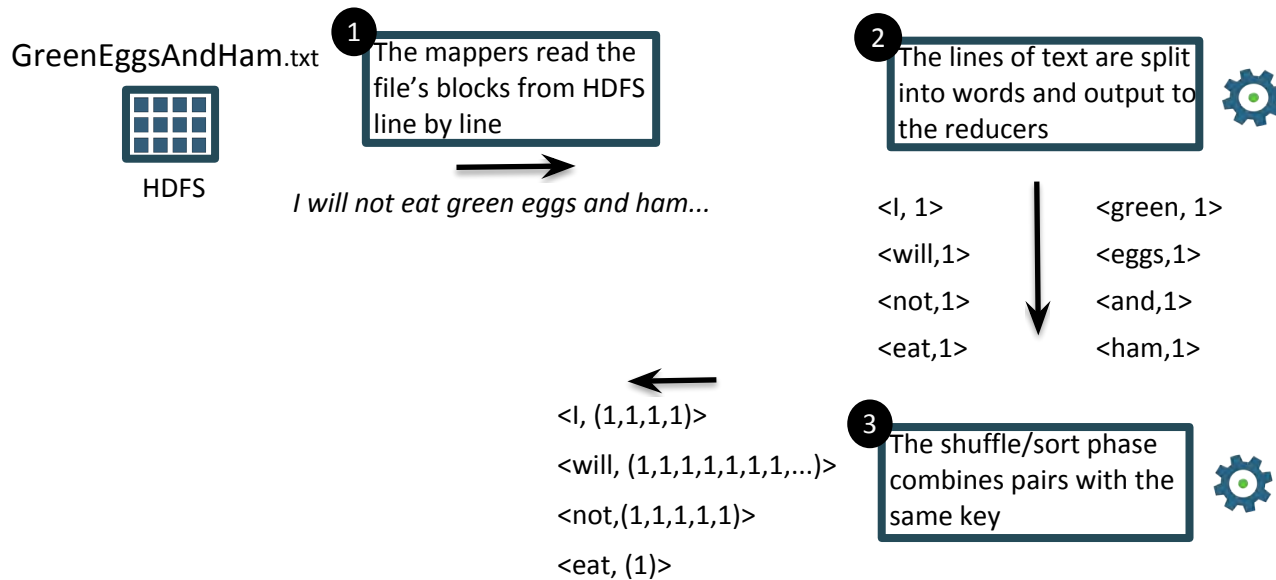
Recap of Word Count



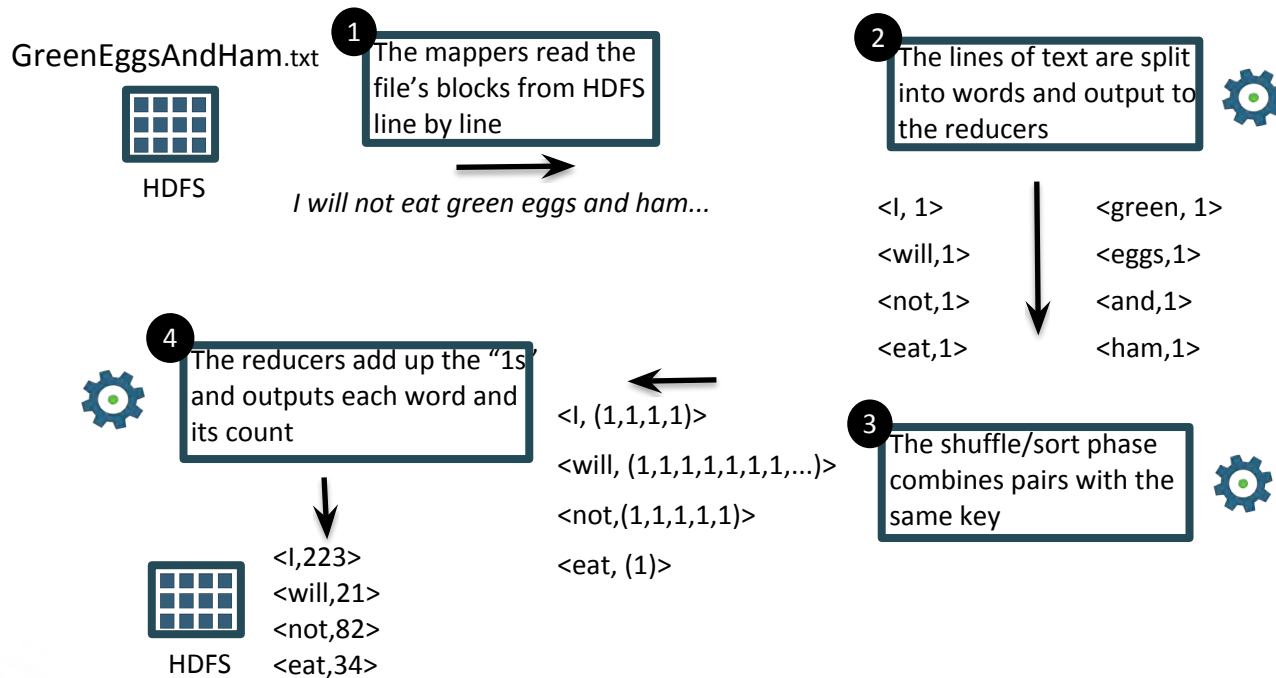
Recap of Word Count



Recap of Word Count



Recap of Word Count



MapReduce Example – Word Count

The mapper

```
public class WordCount {  
  
    public static class TokenizerMapper  
        extends Mapper<Object, Text, Text, IntWritable>{  
  
        private final static IntWritable one = new IntWritable(1);  
        private Text word = new Text();  
  
        public void map(Object key, Text value, Context context  
            ) throws IOException, InterruptedException {  
            StringTokenizer itr = new StringTokenizer(value.toString());  
            while (itr.hasMoreTokens()) {  
                word.set(itr.nextToken());  
                context.write(word, one);  
            }  
        }  
    }  
}
```



MapReduce Example – Word Count

The reducer

```
public static class IntSumReducer
    extends Reducer<Text,IntWritable,Text,IntWritable> {
    private IntWritable result = new IntWritable();

    public void reduce(Text key, Iterable<IntWritable> values,
        Context context
        ) throws IOException, InterruptedException {

        int sum = 0;
        for (IntWritable val : values) {
            sum += val.get();
        }
        result.set(sum);
        context.write(key, result);
    }
}
```



MapReduce Example – Word Count

The main

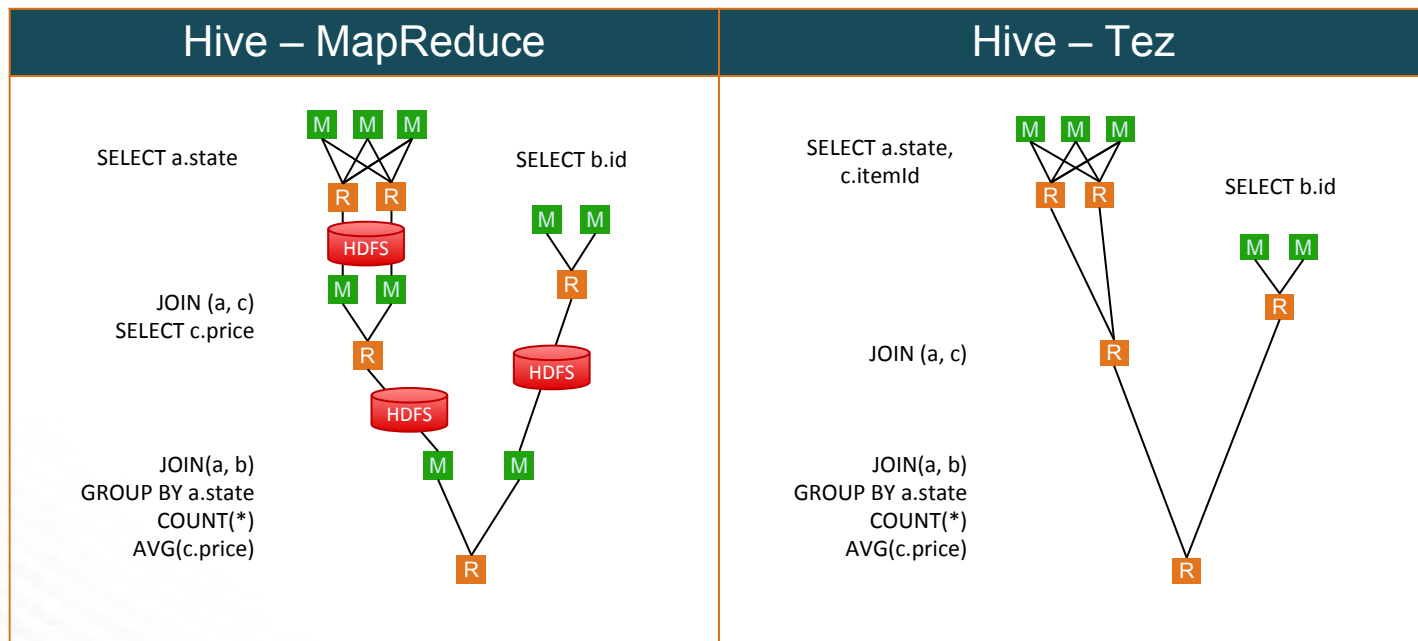
```
public static void main(String[] args) throws Exception {  
    Configuration conf = new Configuration();  
    Job job = Job.getInstance(conf, "word count");  
    job.setJarByClass(WordCount.class);  
    job.setMapperClass(TokenizerMapper.class);  
    job.setCombinerClass(IntSumReducer.class);  
    job.setReducerClass(IntSumReducer.class);  
    job.setOutputKeyClass(Text.class);  
    job.setOutputValueClass(IntWritable.class);  
    FileInputFormat.addInputPath(job, new Path(args[0]));  
    FileOutputFormat.setOutputPath(job, new Path(args[1]));  
    System.exit(job.waitForCompletion(true) ? 0 : 1);  
}
```



'Complex' processing required chaining MapReduce jobs

```
SELECT a.state, COUNT(*), AVG(c.price) FROM a
JOIN b ON (a.id = b.id)
JOIN c ON (a.itemId = c.itemId)
GROUP BY a.state
```

Tez avoids
unnneeded writes to
HDFS



Knowledge Check



Questions

1. What are the two primary phases of the MapReduce framework?
No, this is not a trick question.



Questions

1. What are the two primary phases of the MapReduce framework?
No, this is not a trick question.
2. What dictates the number of Mappers that are run? Same question for the Reducers.



Questions

1. What are the two primary phases of the MapReduce framework?
No, this is not a trick question.
2. What dictates the number of Mappers that are run? Same question for the Reducers.
3. How many input & output KVPs are passed into, and emitted out of, the Mappers? Same question for the Reducers.



Questions

1. What are the two primary phases of the MapReduce framework?
No, this is not a trick question.
2. What dictates the number of Mappers that are run? Same question for the Reducers.
3. How many input & output KVPs are passed into, and emitted out of, the Mappers? Same question for the Reducers.
4. True/False? It is possible to have a Reducer-only job.



Questions

1. What are the two primary phases of the MapReduce framework?
No, this is not a trick question.
2. What dictates the number of Mappers that are run? Same question for the Reducers.
3. How many input & output KVPs are passed into, and emitted out of, the Mappers? Same question for the Reducers.
4. True/False? It is possible to have a Reducer-only job.
5. Why were frameworks like Pig and Hive built on top of MapReduce?
Again, not a trick question...



Summary



Summary

- MapReduce is the foundational framework for processing data at scale because of its ability to break a large problem into any smaller ones
- Mappers read data in the form of KVPs and each call to a Mapper is for a single KVP; it can return 0..m KVPs
- The framework shuffles & sorts the Mappers' outputted KVPs with the guarantee that only one Reducer will be asked to process a given Key's data
- Reducers are given a list of Values for a specific Key; they can return 0..m KVPs
- Due to the fine-grained nature of the framework, many use cases are better suited for higher-order tools

